

Primitive RA Operators

π	Project order does not matter	(1)
σ	Select	(1)
\times	Cartesian Product	(2)
\cup	Union*	(2)
$-$	Difference*	(2)

SQL DML

SELECT DISTINCT WHERE CROSS JOIN UNION EXCEPT (ALL for log)

(SELECT returns a log,
 π returns a set)

* must be union compatible
- same no. of attributes
- domain of attributes must be compatible

Project and Project

$$\pi_{\vec{X}} \pi_{\vec{Y}} R \equiv \pi_{\vec{X}} R$$

You can eliminate any inner project (note that to be well formed $\vec{X} \subseteq \vec{Y}$)

Project and Select

$$\pi_{\vec{X}} \sigma_{P(\vec{Y})} R \equiv \sigma_{P(\vec{Y})} \pi_{\vec{X}} R$$

If $\vec{Y} \subseteq \vec{X}$ you can move a project of \vec{X} inside a select of \vec{Y}

Project and Product

$$\pi_{\vec{X}}(R \times S) \equiv \pi_{\vec{X} \cap \text{Attrs}(R)} R \times \pi_{\vec{X} \cap \text{Attrs}(S)} S$$

Project and Union

$$\pi_{\vec{X}}(R \cup S) \equiv \pi_{\vec{X}} R \cup \pi_{\vec{X}} S$$

Project and Difference

$$\pi_{\vec{X}}(R - S) \supseteq \pi_{\vec{X}} R - \pi_{\vec{X}} S$$

Select and Project

$$\sigma_{P(\vec{X})} \pi_{\vec{X}} R \equiv \pi_{\vec{X}} \sigma_{P(\vec{X})} R$$

Select and Select

$$\sigma_{P_x(\vec{X})} \sigma_{P_y(\vec{Y})} R \equiv \sigma_{P_x(\vec{X}) \wedge P_y(\vec{Y})} R$$

Select and Product

$$\sigma_{P(\vec{X})}(R \times S) \equiv \sigma_{P(\vec{X})} R \times S \iff \vec{X} \subseteq \text{Attrs}(R)$$

If $\vec{X} \subseteq \text{Attrs}(R)$ you can move a select predicate $P(\vec{X})$ onto R.

Select and Union

$$\sigma_{P(\vec{X})}(R \cup S) \equiv \sigma_{P(\vec{X})} R \cup \sigma_{P(\vec{X})} S$$

Select and Difference

$$\sigma_{P(\vec{X})}(R - S) \equiv \sigma_{P(\vec{X})} R - S$$

Product and Union

$$R \times (S \cup T) \equiv (R \times S) \cup (R \times T)$$

Product and Difference

$$R \times (S - T) \equiv (R \times S) - (R \times T)$$

Union and Product

$$R \cup (S \times T) \text{ unable to move } \cup \text{ inside } \times$$

Union and Difference

$$R \cup (S - T) \text{ unable to move } \cup \text{ inside } -$$

Difference and Product

$$R - (S \times T) \text{ unable to move } - \text{ inside } \times$$

Difference and Union

$$R - (S \cup T) \equiv (R - S) - T$$

Derived RA Operators (2)

$$\bowtie \text{ Natural Join } R \bowtie S = \sigma_{R.A_1=S.A_1 \wedge \dots \wedge R.A_m=S.A_m} R \times S$$

$$\bowtie \text{ Semi Join } R \bowtie S = R \bowtie \pi_{A \in \text{Attrs}(R) \cap \text{Attrs}(S)} S$$

$$\cap \text{ Intersection* } R \cap S = R - (R - S)$$

$$\div \text{ Division } R \div S = \pi_{A \in \text{Attrs}(R) - \text{Attrs}(S)} R - \pi_{A \in \text{Attrs}(R) - \text{Attrs}(S)} ((\pi_{A \in \text{Attrs}(R) - \text{Attrs}(S)} R \times S) - R)$$

$$A=B \quad \bowtie \text{ Equi Join } R \bowtie S = \sigma_{R.A_i=S.B_i} R \times S$$

$$\bowtie \text{ Theta Join } R \bowtie S = \sigma_{\phi} R \times S$$

project the extra attributes in R that have each of the attributes of S

SQL DML

NATURAL JOIN

INTERSECT (ALL for log)

SELECT DISTINCT * FROM t AS t;
WHERE NOT EXISTS (SELECT * FROM t WHERE t.x=t.x); } x=y
EXCEPT SELECT * FROM t WHERE t.x=t.x; } x=y

$\pi_{\text{cname}, \text{type}}$ account
'McBrien, P.'
'McBrien, P.'
'Boyd, M.'
'Poulouvassilis, A.'
'Poulouvassilis, A.'
'Bailey, J.'

π_{type} account
'current'
'deposit'

$\pi_{\text{cname}, \text{type}}$ account / π_{type} account
'McBrien, P.'
'Poulouvassilis, A.'

JOIN.. ON

SQL Set Operations

WHERE IN
WHERE no IN (100, 101) listed produced

EXIST WHERE NOT EXIST (SELECT...) produced

ALL WHERE '11-jan-1999' < ALL (SELECT...) produced

SOME WHERE 'deposit' = SOME(SELECT...) produced

SQL Null

nulls are combined in SETS

$x=null, null=null$ gives UNKNOWN

$x \text{ IS NULL}, x \text{ IS NOT NULL}$ gives TRUE/FALSE

$\text{NOT IN}(null, \dots)$ always returns nothing!

$\text{NOT EXIST}(\text{SELECT} \dots = \text{null})$ always returns everything!

SQL JOINS

LEFT (RIGHT) JOIN returns every row in left/right even if no row matches, filled with NULL instead

OUTER JOIN returns LEFT JOIN \cup RIGHT JOIN

FULL OUTER JOIN IN SQL

SQL Extensions

Pattern Matching: WHERE column LIKE pattern ESCAPE escape-char

Cases : , CASE
WHEN ~
THEN ~
;
(ELSE ~) optional
END AS ~,

Modifications : ROUND(n, dp), SUBSTRING(str FROM n, FOR n2),
(must be placed in SELECT clause) UPPER(str), CHAR_LENGTH(str), POSITION(char IN str)

Aggregate : SUM, COUNT, AVG, MIN, MAX excludes NULL

COALESCE(a,b,c) goes down list until non-null

{ for single char
~ for any char
[A-C] for single A to C
[ABC] for single ABC

SQL OLAP

GROUP BY : only one output per group, aggregate functions for non grouped columns. NULL is grouped together

HAVING : similar to WHERE but for aggregates (SELECT.... FROM... WHERE... GROUP BY ... HAVING)

PARTITION: one output per row, $\text{Avg}(x)$ OVER (PARTITION BY y)

ORDER BY ... DESC : default ascending, RANK() OVER (ORDER BY ...) for numbering.

DATALOG Head :- Body

1. If Body then Head.
2. Head is a single pred
3. Body is a conjunction of pred
- a. cannot use same name for intentional & extensional preds.
b. preds start with small letter, vars start with capital letter
c. _ can be used for var that only appears once

Safe negation : A var in a negated \rightarrow pred must also appear in a non-negated pred.

TT : use subset of body vars in head

DI : name two preds in rule body, with same name / compare attribute

σ : use variable more than once / with data value

U : more than one rule def for intentional pred

x : name two preds in rule body

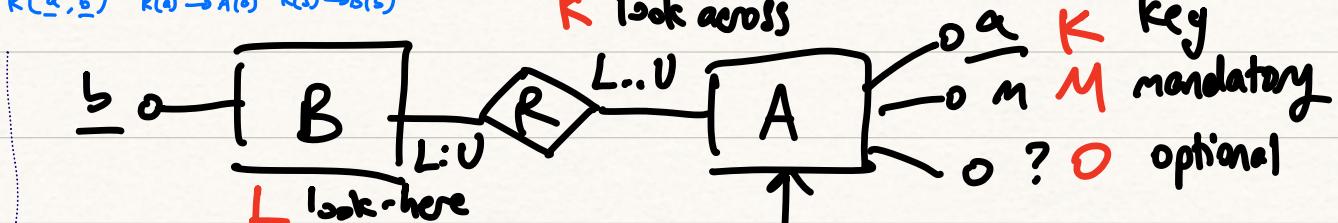
- : negation on pred being subtracted

ER Modelling

$\begin{array}{ll} 1:1-D:N, & A(\dots, b) \xrightarrow{Fk} B(b) \\ D:1-D:N, & A(\dots, b?) \xrightarrow{Fk} B(b) \\ D:N-O;N, & R(\underline{a}, \underline{b}) \xrightarrow{Fk} A(a) \quad R(\underline{c}) \xrightarrow{Fk} B(b) \end{array}$

B appears L-U times in A
K look across

$A(\underline{a}, m, o?)$

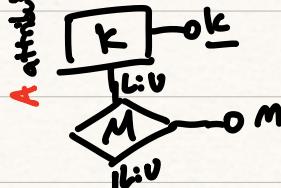


B appears L-U times in A

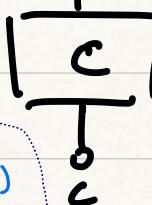
S subset/isa hierarchy

$c(\underline{a}, c)$ $c(a) \xrightarrow{Fk} A(a)$

2 nested relationship



$M(E, h, i, j, m)$ $M(h) \xrightarrow{Fk} K(h)$
 $M(h, i, j) \xrightarrow{Fk} N(h, i, j)$



$I(i)$
 $IM(I, \underline{a}) \xrightarrow{Fk} I(\underline{a})$
 $IS(I, \underline{i}) \xrightarrow{Fk} I(\underline{i})$ } * and + some
 $IS(I) \xrightarrow{Fk} I(\underline{i})$ } all key

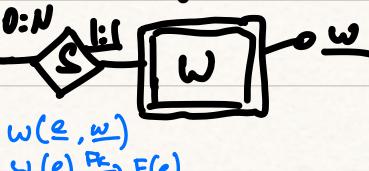


OM + V multi-valued
optional + multivalued

many-many, $N(h, i, j)$ $3 \times Fk$ all key
one-many from H, $H(b, i, j)$ $2 \times Fk$

at least 1 1:1 attribute & key

W weak entity



disjointness / generalization hierarchy

D

Keys & FDs

Super-key X of R: X determines all other attributes of R

Generate minimal keys:
 1. Generate minimal cover
 2. Find keys where $K^+ = \text{Attr}(R)$

(non-unique) minimal-key X of R: not possible to remove attributes to form another super-key.

Closure S^+ : all FDs that can be inferred from S. $S^+ = T^+$ if $S^+ = T^+$

minimal cover S_C : $S_C^+ = S^+$ and not possible to remove another FD

Generate minimal cover:

Normalisation
 1st Normal Form (1NF) Every attribute depends on key.

1. rewrite S to have single attribute on RHS of each FD
2. For each FD $X \rightarrow A$, if $(X-B) \rightarrow B$, remove B from X
3. For each FD $X \rightarrow A$, compute X^+ without $X \rightarrow A$. If $A \subseteq X^+$, remove it

Prime Attributes Attributes A of R that are part of a minimal candidate key X of R

3rd Normal Form (3NF) For every non-trivial FD $X \rightarrow A$, either X is a super-key or A is prime

Boyce-Codd Normal Form (BCNF) For every non-trivial FD $X \rightarrow A$, X is a super key.

Lossless-join decomposition Decompose R into R_1, \dots, R_n such that $\text{Attr}(R_1) \cup \dots \cup \text{Attr}(R_n) = \text{Attr}(R)$

FD-preserving decomposition Lossless decomposition preserve FDs if $S^+ = (S_a \cup S_b)^+$ may need to add more tables on top of lossless

Generating 3BCNF: (1) Find $\text{FD}(X \rightarrow A) \subseteq S$ which violates not possible to preserve FD NF (X not superkey and A non-prime 3NF BCNF)

(2) Decompose R into $R_a(\text{Attr}(R)-A)$ and $R_b(XA)$ ($X \subseteq R_a$ & $(X \rightarrow A) \subseteq R_b$)

(3) Project S onto new relations and repeat from (1)

Serialisability & Recoverability

Serialisability concurrent execution same end result as some serial execution of those transaction

Recoverability No transaction commits depending on uncommitted data

(1) Lost Update not serialisable
recoverable $r_i[0] \leftarrow w_1[0] \leftarrow w_2[0]$ reading update leads to no change

(2) Inconsistent Analysis not serialisable
recoverable $r_i[0_a] \leftarrow w_2[0_a], w_2[0_b] \leftarrow r_i[0_b]$ reading halfway through update

(3) Dirty Read serialisable
may not recoverable $w_1[0] \leftarrow r_2[0] \leftarrow e_1$ read uncommitted data

(4) Dirty Write may not serialisable
recoverable $w_1[0] \leftarrow w_2[0] \leftarrow e_1$ overwrite uncommitted data

Serialisability conflicts occur between $(r_x[0] \text{ and } w_y[0])$ or $(w_x[0] \text{ and } w_y[0])$ in H

H is conflict serialisable (CSR) if the serialisation graph of t is acyclic. e.g. lost update (1) lost update inconsistent analysis (2) inconsistent analysis cycle

Recoverable (RC) if no transactions commit before another transaction it read from commits

Avoids Cascading Aborts (ACA) if no reads from non-committed transaction no dirty read ST C ACA C RC

Strict (ST) if no read from non-committed transaction + no write over non-committed transaction no dirty read no dirty write

Concurrency

Two-Phase Locking (2PL)

(1) growing phase + shrinking phase \rightarrow only one peak / cannot unlock then lock

(2) refuse $rl_i[0]$ if $wl_j[0]$ is held } deadlock transaction

(3) refuse $wl_i[0]$ if $rl_j[0]$ or $wl_j[0]$ is held } put on wait-for-graph (WFG)

If WFG has no cycle \rightarrow no deadlock

If WFG has cycle \rightarrow deadlock \rightarrow rollback one of the transactions (abort)