COMP50008 Probability & Statistics
Imperial College London

Boxuan Tang

Spring 2023

Contents

1 Probability

1.1 Events

Definition 1.1 (Sample Spaces) *A set that describes the rage of possible outcomes of a random experiment*

Definition 1.2 (Events) *An event E is any subset of the sample space $E \subseteq S$, a collection of some of its possible outcomes. The singleton subsets of S are the **elementary** events of S. The events are **mutually exclusive** if $\forall i, j, E_i \cap E_j = \emptyset$ and at most one of the events can occur.*

Definition 1.3 (σ-algebra) *The set of subsets $\mathcal{F}$ must be*

- *nonempty ($S \in \mathcal{F}$)*
- *closed under complement ($E \in \mathcal{F} \to \overline{E} \in F$)*
- *closed under countable union ($E_1, \cdots \in \mathcal{F} \to \bigcup_i E_i \in \mathcal{F}$)*

Definition 1.4 (Axioms of Probability) *A **probability measure** on the pair $(S, \mathcal{F})$ is a mapping $P : \mathcal{F} \to [0, 1]$ satisfying the following axioms for all subsets of S*

- $\forall E \in \mathcal{F}, 0 \leq P(E) \leq 1$
- $P(S) = 1$
- *countably additive meaning for disjoint subsets $E_1 \cdots \in \mathcal{F}, P(\bigcup_i E_i) = \Sigma_i P(E_i)$*

Definition 1.5 (Independent Events) *A set of events $\{E_1, E_2, \dots\}$ is independent iff for any finite subset $\{E_{i_1} \dots E_{i_n}\}$, $P(\bigcap_{j=1}^{n} E_{i_j}) = \Pi_{j=1}^{n} P(E_{i_j})$. Two events E and F are independent iff $P(E \cap F) = P(E) \cap P(F)$. If E and F are independent, $\overline{E}$ and F are also independent.*
$P(E \cup F) = P(E) + P(F) - P(E \cap F)$

Definition 1.6 (Conditional Probability) $P(E \mid F) = \frac{P(E \cap F)}{P(F)}$. *The events E_1 and E_2 are **conditionally independent** given F iff $P(E_1 \cap E_2 \mid F) = P(E_1 \mid F) P(E_2 \mid F)$.*

Theorem 1.7 (Bayes Theorem) $P(E \mid F) = \frac{P(E) P(F \mid E)}{P(F)}$.

Theorem 1.8 (Partition Rule) $P(E) = \Sigma_i P(E \mid F_i) P(F_i) = P(E \mid F) P(F) + P(E \mid \overline{F}) P(\overline{F})$

1.2 Random Variables

Definition 1.9 (Random Variable) *A random variable X is a mapping from the sample space to real numbers, $X : S \to \mathbb{R}$. The **support** of X is the image of S under X, $supp(X) = X(S) = \{x \in \mathbb{R} \mid \exists s \in S, X(s) = x\}$*

Definition 1.10 (Cumulative Distribution Function) *The cumulative distribution function $F_x(x)$ is the probability that X takes a value less than or equal to x, $F_x(x) = P_x(X \leq x)$. For it to be valid:*

- ***(Monotonic)*** $\forall x_1, x_2 \in \mathbb{R}, x_1 < x_2 \to F_x(x_1) \leq F_x(x_2)$
- $F_x(-\infty) = 0, F_x(\infty) = 1$
- *F_x is right-continuous*

$P_x(a < X \leq b) = F_x(b) - F_x(a)$

Definition 1.11 (Expectation) ***Discrete:*** $\mu = E(X) = \Sigma_x x p(x)$.
Continuous: $\mu = E(x) = \int_{-\infty}^{\infty} x f_X(x) dx$ *or generally* $E(g(X)) = \int_{-\infty}^{\infty} g(x) f_X(x) dx$
Linearity of Expectation: $E(aX + b) = aE(x) + b$.

Definition 1.12 (Variance) $Var(X) = E[(X - E(X))^2] = E(X^2) - (E(X))^2$. $Var(aX + b) = a^2 Var(X)$
Standard Deviation $\sigma = \sqrt{Var(X)}$.

1.3 Discrete Random Variables

Definition 1.13 (Probability Mass Function) *For a discrete random variable X, the probability mass function is* $p(x) = P_x(X = x)$ *where* $0 \leq p(x) \leq 1$ *and* $\Sigma_{x\in\mathcal{X}} p(x) = 1$. *Also* $p(x_i) = F(x_i) - F(x_{i-1})$ *and* $F(x_i) = \Sigma_{j=1}^{i} p(x_j)$

Definition 1.14 (Skewness) $\gamma_1 = \frac{E[(X-\mu)^3]}{\sigma^3}$

Definition 1.15 (Sum of Random Variables) *Let* $X_1 \dots X_n$ *be random variables and* $S_n = \Sigma_{i=1}^{n} X_i$. $E(S_n) = \Sigma_{i=1}^{n} E(X_i)$ *and* $E(\frac{S_n}{n}) = \frac{\Sigma_{i=1}^{n} E(X_i)}{n}$. $Var(S_n) = \Sigma_{i=1}^{n} Var(X_i)$ *and* $Var(\frac{S_n}{n}) = \frac{\Sigma_{i=1}^{n} Var(X_i)}{n^2}$.
If they are independent and identically distributed where $E(X_i) = \mu_x$. *and* $Var(X_i) = \sigma_X^2$, $E(\frac{S_n}{n}) = \mu_x$ *and* $Var(\frac{S_n}{n}) = \frac{\sigma_X^2}{n}$.

Distribution	rv	pmf	μ	σ^2	γ_1	MLE
Bernoulli(p)	$0, 1$	$p^x(1-p)^x$	p	p(1-p)		
Binomial(n,p)	$0 \dots n$	$\binom{n}{x} p^x(1-p)^x$	np	np(1-p)	$\frac{1-2p}{\sqrt{np(1-p)}}$	$\overline{x}$
Geometric(p)	$1 \dots$	$p(1-p)^{x-1}$	$\frac{1}{p}$	$\frac{1-p}{p^2}$	$\frac{2-p}{\sqrt{1-p}}$	$\frac{1}{\overline{x}}$
Poisson(λ)	$0 \dots$	$\frac{e^{-\lambda}\lambda^x}{x!}$	λ	λ	$\frac{1}{\sqrt{\lambda}}$	$\overline{x}$
Uniform(1,n)	$1 \dots n$	$\frac{1}{n}$	$\frac{n+1}{2}$	$\frac{n^2-1}{12}$		

when p is small and n is large, Binomial(n,p) $\sim$ Poisson(np)

1.4 Continuous Random Variables

Definition 1.16 (Probability Density Function) *The probability density function* f_X *of a* ***continuous*** *X is such that* $F_X(x) = \int_{u--\infty}^{x} f_X(u)du$ *where* $\forall x \in \mathbb{R}. f_X(x) \geq 0$ *and* $\int_{-\infty}^{\infty} f_X(x)dx = 1$. *So* $f_X(x) = \frac{d}{dx} F_X(x)$. *Also* $P_X(a < X \leq b) = \int_a^b f_X(x)dx$.

Definition 1.17 (Quantile) α*-quantile* $Q_X(\alpha), 0 \leq \alpha \leq 1$ *is the least number satisfying* $P(X \leq Q_X(\alpha)) = \alpha$, $Q_X(\alpha) = F_X^{-1}(\alpha)$. *The median is the* $\frac{1}{2}$*-quantile and the* k^{th} *percentile is the* $\frac{k}{100}$*-quantile.*

Distribution	pdf	cdf	validity	μ	σ^2	MLE
U(a,b)	$\frac{1}{b-a}$	$\frac{x-a}{b-a}$	$a < x < b$	$\frac{a+b}{2}$	$\frac{(b-a)^2}{12}$	
Exp(λ)	$\lambda e^{-\lambda x}$	$1 - e^{-\lambda x}$	$x \geq 0$	λ	λ^2	$\overline{x}$
N(μ, σ^2)	$\frac{1}{\sigma\sqrt{2\pi}} e^{-\frac{(x-\mu)^2}{2\sigma^2}}$	$\frac{1}{\sigma\sqrt{2\pi}} \int_{-\infty}^{x} e^{-\frac{(x-\mu)^2}{2\sigma^2}}$	all	μ	σ^2	$(\overline{x}, S_{n-1}^2)$
Lognormal(μ, σ^2)	$\frac{1}{\sigma x\sqrt{2\pi}} e^{-\frac{(\ln x-\mu)^2}{2\sigma^2}}$		+ve	$e^{\mu+\frac{\sigma^2}{2}}$	$e^{\sigma^2} - 1$	

If events in a random process $\sim Poisson(\lambda)$ then the time between consecutive events $\sim Exp(\lambda)$.
$X \sim N(\mu, \sigma^2) \implies \frac{X-\mu}{\sigma} \sim \Phi$ and $\Phi(1.96) \approx 0.975, \Phi(2.58) \approx 0.995$.

Definition 1.18 (Moment Generating Function) $E[X^n] = \frac{d^n M_X(t)}{dt^n} \mid_{t=0}$ *where*

$$M_X(t) = E(e^{tX}) = \int_{-\infty}^{\infty} e^{tX} f_X(x)dx$$

Definition 1.19 (Sum of Independent Random Variables) $M_{S_n}(t) = \Pi_{j=1}^{n} M_{X_j}(t)$ *and* $E[\Pi_{i=1}^{n} Z_i] = \Pi_{i=1}^{n} E[Z_i]$. *For 2 variables,* $M_{Z_1+Z_2}(t) = M_{Z_1}(t)M_{Z_2}(t)$ *and* $E[Z_1 Z_2] = E[Z_1]E[Z_2]$.

Theorem 1.20 (Central Limit Theorem) *Let* $X_1 \dots X_n$ *be independent and identically distributed random variables from any distribution with mean* μ *and* ***finite*** *variance* σ^2. *Then* $\lim_{n\to\infty} \frac{S_n - n_u}{\sqrt{n}\sigma} = \frac{\overline{X}-\mu}{\sigma/\sqrt{n}} \sim \phi$

1.5 Joint Random Variables

Definition 1.21 (Joint Cumulative Distribution Function) $F(x, y) = P_Z(X \leq x, Y \leq y)$
so $F_X(x) = F(x, \infty)$ *and* $F_Y(y) = F(\infty, y)$

- $\forall x, y \in \mathbb{R}, 0 \leq F(x, y) \leq 1$
- ***Monotonicity*** $x_1 < x_2 \implies F(x_1, y_1) \leq F(x_2, y_1)$ *and* $y_1 < y_2 \implies F(x_1, y_1) \leq F(x_1, y_2)$

- $\forall x, y \in \mathbb{R}. F(x, -\infty) = F(-\infty, y) = 0$ *and* $F(\infty, \infty) = 1$

$P_Z(x_1 < X \leq x_2, y_1 < Y \leq y_2) = F(x_2, y_2) - F(x_1, y_2) - F(x_2, y_1) + F(x_1, y_1)$

Definition 1.22 (Joint Probability Mass Function) $p(x, y) = P_Z(X = x, Y = y)$ *where* $\forall x, y \in \mathbb{R}, 0 \leq p(x, y) \leq 1$ *and* $\Sigma_y \Sigma_x p(x, y) = 1$. *So* $p_X(x) = \Sigma_y p(x, y)$ *and* $p_Y(y) = \Sigma_x p(x, y)$.

Definition 1.23 (Joint Probability Density Function) $F(x, y) = \int_{t=-\infty}^{y} \int_{s=-\infty}^{x} f(s, t) ds dt$ *where* $\forall x, y \in \mathbb{R}, f(x, y) \geq 0$ *and* $\int_{y=-\infty}^{\infty} \int_{x=-\infty}^{\infty} f(x, y) dx dy = 1$. *So* $f(x, y) = \frac{\partial^2}{\partial x \partial y} F(x, y)$. *Also* $f_X(x) = \int_{y=-\infty}^{\infty} f(x, y) dy$ *and* $f_Y(y) = \int_{x=-\infty}^{\infty} f(x, y) dx$ ***(Marginal Density Functions)***

Definition 1.24 (Independence) ***Discrete****:* $p(x, y) = p_X(x) p_Y(y)$. ***Continuous****:* $f(x, y) = f_X(x) f_Y(y)$.

Definition 1.25 (Partition Rule) ***Discrete****:* $p_X(x) = \Sigma_y p_{X|Y}(x \mid y) p_Y(y)$.
Continuous*:* $f_X(x) = \int_{y-\infty}^{\infty} f_{X|Y}(x \mid y) f_Y(y) dy$.

Definition 1.26 (Expectation) ***Discrete****:* $E(g(X, Y)) = \Sigma_y \Sigma_x g(x, y) p(x, y)$.
Continuous*:* $E(g(X, Y)) = \int_{y=-\infty}^{\infty} g(x, y) f(x, y) dx dy$.
If $g(X, Y) = g_1(X) + g_2(Y)$ *then* $E(g(X, Y)) = E_X(g_1(X)) + E_Y(g_2(Y))$.
If $g(X, Y) = g_1(X) g_2(Y)$ *and* X *and* Y *are independent, then* $E(g(X, Y)) = E_X(g_1(X)) E_Y(g_2(Y))$.

Definition 1.27 (Covariance) $\sigma_{XY} = E[XY] - \mu_X \mu_Y$. *For independent rvs,* $\sigma_{XY} = 0$.

Definition 1.28 (Correlation) $\rho_{XY} = \frac{\sigma_{XY}}{\sigma_X \sigma_Y}$.

Definition 1.29 (Conditional Expectation) ***Discrete****:* $E_{Y|X}(Y \mid x) = \Sigma_y y p_{Y|X}(y \mid x)$.
Continuous*:* $E_{Y|X}(Y \mid x) = \int_{y=-\infty}^{\infty} y f_{Y|X}(y \mid x) dy$. *Also,* $E_Y(Y) = E_X(E_{Y|X}(Y \mid X))$.
Tower Rule*:* $E(Y) = E_{X_n}(E_{X_{n-1}}(\dots E_{X_1}(E_Y(Y \mid X_1 \dots X_n) \mid X_2 \dots X_n) \cdots \mid X_n))$.

Definition 1.30 (Discrete Time Markov Chain) $P(X_n = j) = (\pi_0 R^n)_j$ *where* $P(X_0 = i) = \pi_{0i}$ *for the horizontal initial probability vector* π_0 *and* $r_{ij} = P(X_{n+1} = j \mid X_n = i)$ *for the transition matrix* R.
Since $\pi_\infty R = \pi_\infty$, R *has an eigenvalue of 1 with the eigenvector* π_∞.

2 Statistics

Definition 2.1 (Bias) *The bias of an estimator* T *for a parameter* θ *is* $bias(T) = E[T \mid \theta] - \theta$.
If the estimator has zero bias we say the estimator is unbiased.

Definition 2.2 (Variance) $S^2 = \frac{1}{n} \Sigma_{i=1}^{n} (x_i - \overline{x})^2$. ***Bias-corrected Sample Variance****:*

$$S_{n-1}^2 = \frac{1}{n-1} \Sigma_{i=1}^{n} (x_i - \overline{x})^2 = \frac{1}{n-1} (\Sigma_i x_i^2 - n\overline{x}^2) = \frac{n}{n-1} S^2$$

Definition 2.3 (Efficiency of Estimators) *For 2 unbiased estimators* $\hat{\Theta}$ *and* $\tilde{\Theta}$, $\hat{\Theta}$ *is more efficient than* $\tilde{\Theta}$ *if* $\forall \theta. Var_{\hat{\Theta}|\theta}(\hat{\Theta} \mid \theta) \leq Var_{\tilde{\Theta}|\theta}(\tilde{\Theta} \mid \theta)$ *and* $\exists \theta. Var_{\hat{\Theta}|\theta}(\hat{\Theta} \mid \theta) < Var_{\tilde{\Theta}|\theta}(\tilde{\Theta} \mid \theta)$

Definition 2.4 (Consistency of Estimators) $\hat{\Theta}$ *is a consistent estimator for* θ *if* $\forall \epsilon > 0. P(\left|\hat{\Theta} - \theta\right| > \epsilon) \to 0$ *as* $n \to \infty$. *If* $\hat{\Theta}$ *is unbiased then* $\lim_{n \to \infty} Var(\hat{\Theta}) = 0 \implies \hat{\Theta}$ *is consistent.*

Definition 2.5 (Confidence Interval) *For known population variance* σ^2, *the 100(1-*α*)% CI for* μ *is*

$$[\overline{x} - z_{1-\frac{\alpha}{2}} \frac{\sigma}{\sqrt{n}}, \overline{x} + z_{1-\frac{\alpha}{2}} \frac{\sigma}{\sqrt{n}}]$$

State CLT if used*. For 2-tailed 95% CI,* $z = 1.96$. *Otherwise, the 100(1-*α*)% confidence interval for* μ *is*

$$[\overline{x} - t_{n-1,1-\frac{\alpha}{2}} \frac{S_{n-1}}{\sqrt{n}}, \overline{x} + t_{n-1,1-\frac{\alpha}{2}} \frac{S_{n-1}}{\sqrt{n}}]$$

Definition 2.6 (Hypothesis Testing) *Identify the rejection region R of T under the assumption H_0 is true, $P(T \in R \mid H_0) = \alpha$.*
State CLT if used. State assume H_0 to be true. State at what α-level.
Testing if $\overline{X} = \mu_0$: For known population variance σ^2, $Z = \frac{\overline{X}-\mu_0}{\sigma/\sqrt{n}} \sim \Theta$. Otherwise, $T = \frac{\overline{X}-\mu_0}{s_{n-1}/\sqrt{n}} \sim t_{n-1}$.
Testing if $\mu_X = \mu_Y$. For known population variance σ^2, $Z = \frac{\overline{X}-\overline{Y}}{\sqrt{\sigma_X^2/n_1+\sigma_Y^2/n_2}} \sim \Theta$. Otherwise,

$$T = \frac{\overline{X}-\overline{Y}}{S_{n_1+n_2-2}\sqrt{1/n_1+1/n_2}} \sim t_{n_1+n_2-2} \qquad S_{n_1+n_2-2} = \frac{n_1-1}{n_1+n_2-2}S_{n_1-1}^2 + \frac{n_2-1}{n_1+n_2-2}S_{n_2-1}^2$$

***Type I** error is rejecting H_0 when it is true. **Type II** error is not rejecting H_0 when H_1 is true.*
***Power of Test**: $P(T \in R \mid H_1)$, high probability of rejecting H_0 when H_1 is true.*

Definition 2.7 (Chi-Squared Test)

$$X_{k-p-1}^2 = \Sigma_{i=1}^k \frac{(O_i - E_i)^2}{E_i}$$

Define H_0 and H_1. If $x^2 < X_{k-p-1,1-\alpha}^2$, where k is the number of values X can take and p is the number of parameters being estimated, we do not reject the H_0 at the α significance level.
***Independence Test**: For $k \times l$ observed, expected value is $\hat{n}_{ij} = \frac{n_{i\bullet} \times n_{\bullet j}}{n}$ with $(k-1)(l-1)$ degrees of freedom.*

Definition 2.8 (Likelihood) *$L(\theta) = P(X \mid \theta)$ or $l(\theta) = \log P(X \mid \theta)$.*
Maximum Likelihood Estimate $= argmax_\theta L(\theta \mid x)$:

1. *Write down $L(\theta) = \Pi_{i=1}^n f(x_i \mid \theta)$ which is the product of the n pdf/pmf viewed as a function of θ.*
2. *Take the natural log of the likelihood to get $l(\theta)$.*
3. *Find the value of θ where $l(\theta)$ is maximised by solving $\frac{\partial}{\partial\theta} l(\theta) = 0$*
4. ***Check** that the estimate in step 3 is a maximum by checking that $\frac{\partial^2}{\partial\theta^2} l(\theta) < 0$*

Definition 2.9 (Posterior) *Posterior $=$ Likelihood$\times$Prior$\times \frac{1}{Evidence}$ or $P(\theta \mid X) = P(X \mid \theta) \times P(\theta) \times \frac{1}{P(X)}$*
***Bayesian Estimate**: $\hat{\theta}_B$ is the mean of the new distribution*
***Maximum A Posteriori Estimate**: $\hat{\theta}_{MAP} = argmax_\theta[\Pi_{i=1}^n P(X = x_i \mid \theta) \times P(\theta)]$*

Definition 2.10 (Beta Prior) *Used for Bernoulli, Binomial and Geometric distributions.*

$$Beta(\theta;\alpha,\beta) = \frac{\theta^{\alpha-1}(1-\theta)^{\beta-1}}{\int_0^1 \theta^{\alpha-1}(1-\alpha)^{\beta-1}d\theta} \quad \max = \frac{\alpha-1}{\alpha+\beta-2} \quad \mu = \frac{\alpha}{\alpha+\beta} \quad \sigma^2 = \frac{\alpha\beta}{(\alpha+\beta)^2(\alpha+\beta+1)}$$

Given a $Beta(\theta;\alpha,\beta)$ prior, with sample size n and mean $\overline{x}$, the posterior is $Beta(\theta;\alpha+n\overline{x},\beta+n(1-\overline{x}))$

Definition 2.11 (Gamma Prior) *Used for Poisson and Exponential distributions.*

$$Gamma(\theta;\alpha,\beta) = \frac{\theta^{a-1}e^{-b\theta}}{\int_0^\infty \theta^{a-1}e^{-b\theta}d\theta} \quad \max = \frac{\alpha-1}{\beta} \quad \mu = \frac{\alpha}{\beta}$$

Given a $Gamma(\theta;\alpha,\beta)$ prior, with sample size n and mean $\overline{x}$, the posterior is $B(\theta;\alpha+n\overline{x},\beta+n))$

Definition 2.12 (Normal Prior) *Used for Normal distributions.* $\max = \mu, \mu = \mu$
Given a $N(\mu_0,\sigma_0^2)$ prior, with sample size n of $N(\mu,\sigma_x^2)$ where sample mean is $\overline{\mu}$ and the variance is known, the posterior is $N(\mu_1,\sigma_1^2)$ where

$$\frac{1}{\sigma_1^2} = \frac{1}{\sigma_0^2} + \frac{n}{\sigma_x^2} \qquad \frac{\mu_1}{\sigma_1^2} = \frac{\mu_0}{\sigma_0^2} + \frac{n\overline{x}}{\sigma_x^2}$$